PHASE 8 / SAME-CONTEXT FULL VS APPEND / 16K HISTORY

Extra delay spread across 128 output tokens

Single RTX 4090 | Qwen2.5-7B | 16K history + 1K new input | Independent background KV | 5 repeats

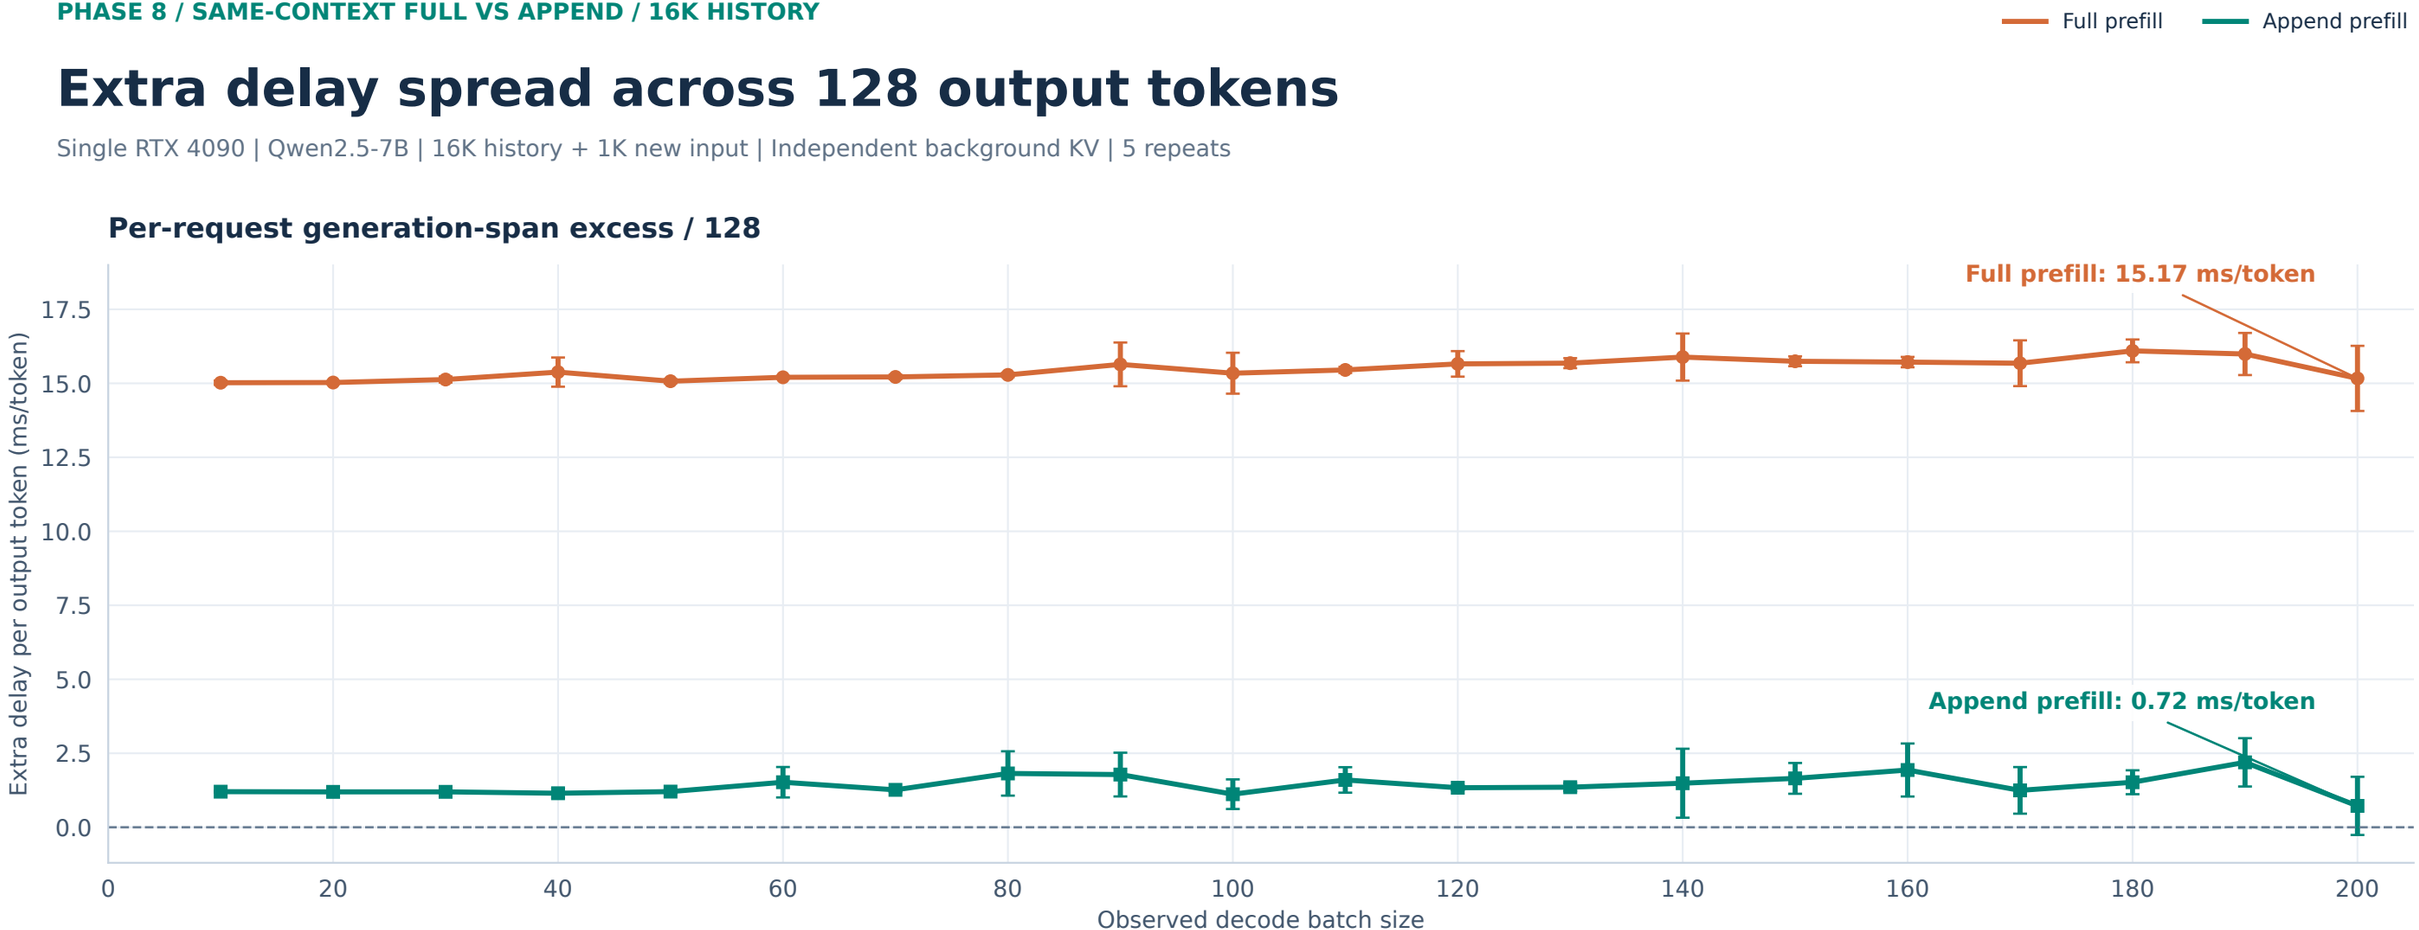

Per trial: [(last - first token arrival) with prefill - matched decode-only span] / 128 = 127/128 x (TPOT - baseline TPOT).

Average background requests first, then five paired repeats; error bars = sample SD. Includes all 127 gaps and signed fluctuations; excludes TTFT.

This is an amortized accounting metric, not a uniform delay at every token. No division by batch size; batch 0 has no background output.