PHASE 8 / SAME-CONTEXT FULL VS APPEND / 16K HISTORY

The stall appears at injection

Single RTX 4090 | Qwen2.5-7B | Independent 128-token decode histories | 16K history + 1K new input | 5 repeats

Offset 0 = actual mixed scheduler step. Matched controls use the same repeat, request index and output-token ordinal. Shading = sample SD, n=5.
Engine iteration includes scheduling/host overhead. Client intervals are singleton token-ID SSE arrival gaps; neither is pure GPU kernel time.
Independent 128-token histories. Full recomputes 17,408 tokens; append reuses 16,384 and computes 1,024. All spikes are retained.