PHASE 8 / SAME-CONTEXT FULL VS APPEND / 16K HISTORY

Full recomputation versus cached append

Single RTX 4090 | Qwen2.5-7B | Independent 128-token decode histories | 16K history + 1K new input | 5 repeats

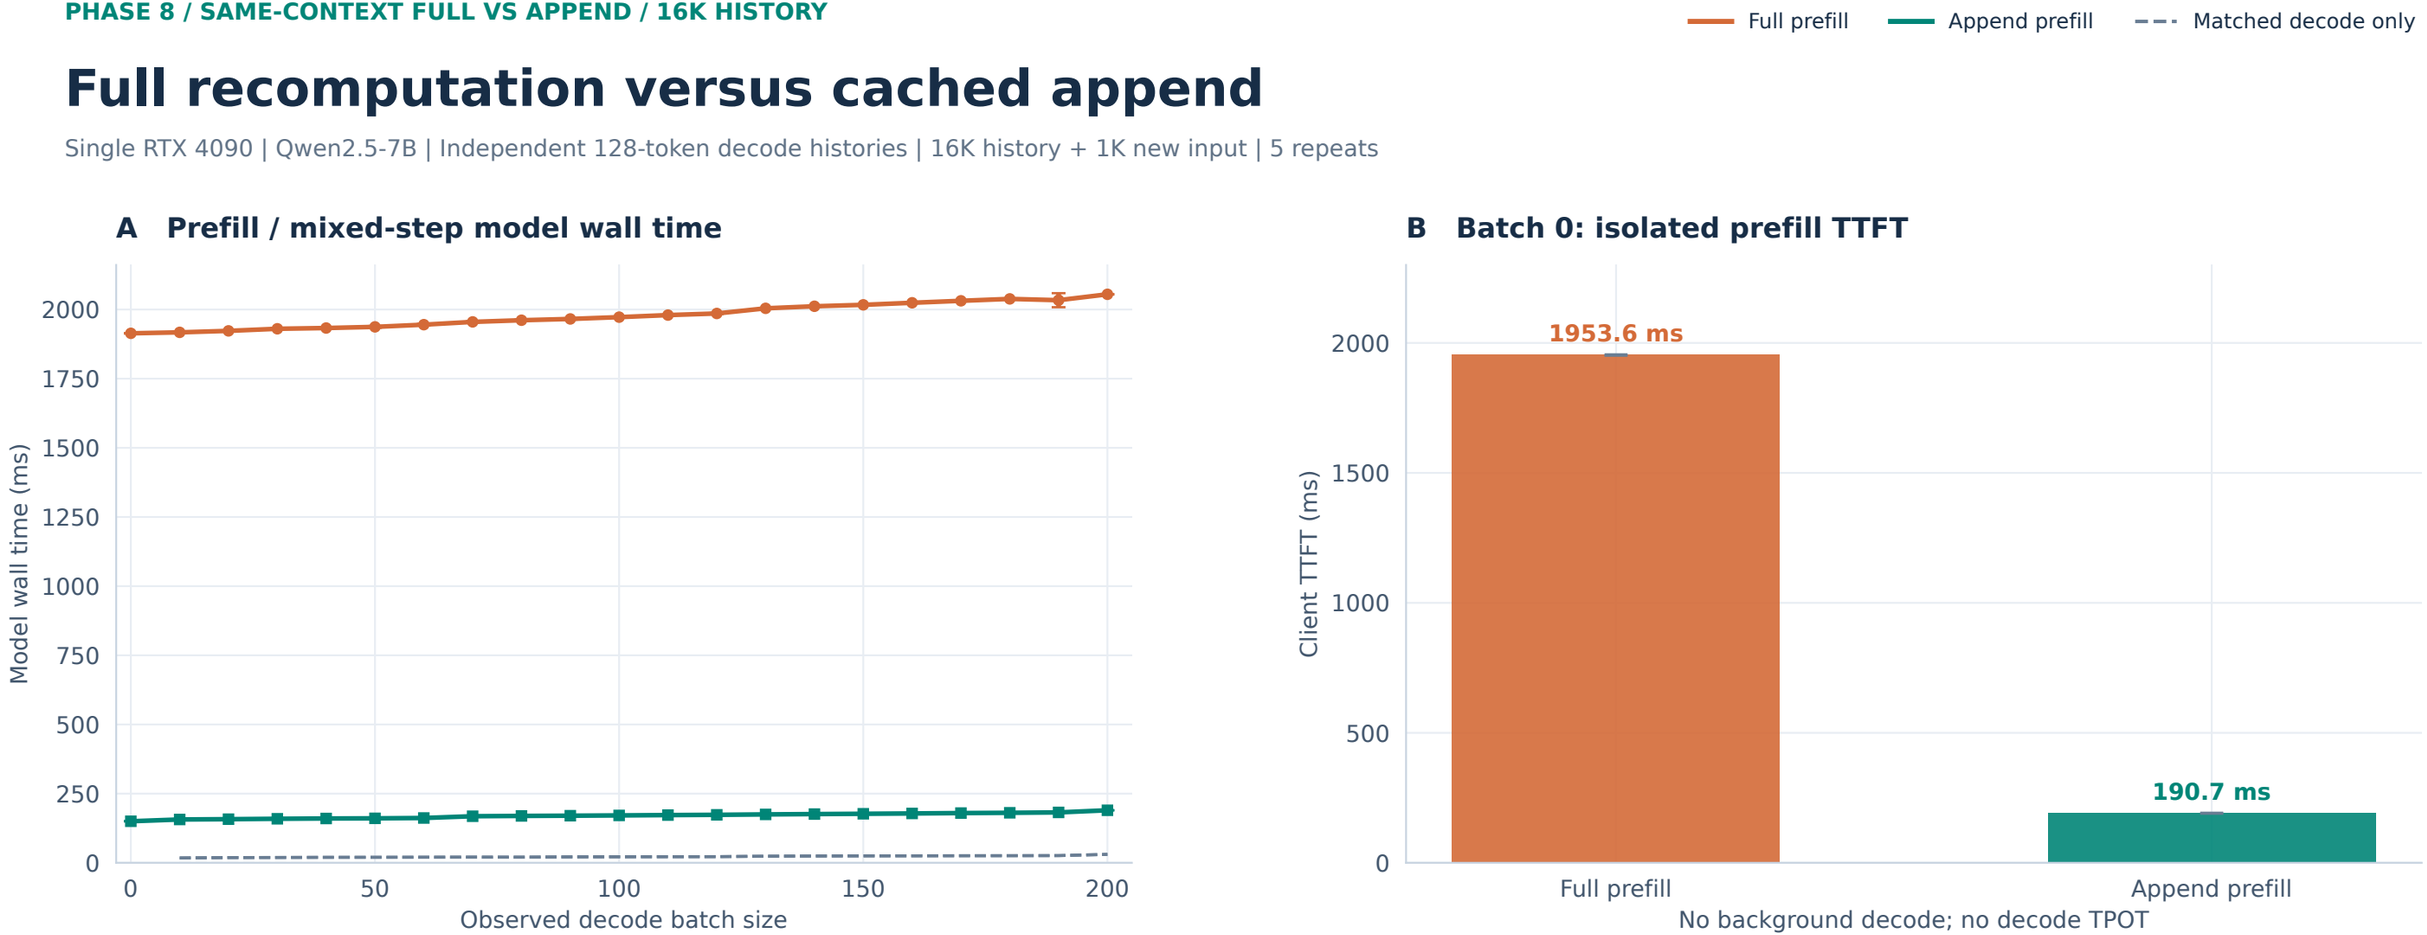

Full and append use exactly the same 16,384 history + 1,024 new tokens. Cold full uses a separate cache salt; append uses the warmed salt.

At batch 0, panel A measures a standalone prefill step. At positive batches it measures prefill mixed with B decode sequences; controls have no prefill.

Means and sample SD across five repeats. Model wall time includes engine/host work; TTFT also includes request handling. Not original-paper numerical replication.